

```

sim_done = 1
sim_done = 4
sim_done = 3
sim_done = 2
sim_done = 5
sim_done = 6
PASS! SimTime = 1.92208e+10
=Worker Thread=
    Throughput:          148207
    MaxThr:              148203
    PerThdThr:           148207
    run_time:            115.297
(19.2161,19.216,19.2163,19.2167,19.2156,19.2161,0,)
    log_bytes:           0 (0,0,0,0,0,0,0,)
    log_dep_size:        0 (0,0,0,0,0,0,0,)
    log_total_size:      0 (0,0,0,0,0,0,0,)
    latency:             253.288
(42.2184,42.2165,42.217,42.2109,42.2118,42.2131,0,)

    num_commits:         2.84797e+06
(500000,474687,499111,401651,498028,474496,0,)
    num_async_commits:   0 (0,0,0,0,0,0,0,)
    num_aborts:          59 (13,5,24,1,13,3,0,)
    num_aborts_logging:  0 (0,0,0,0,0,0,0,)
    num_log_records:     0 (0,0,0,0,0,0,0,)
    log_data:            0 (0,0,0,0,0,0,0,)
    num_latency_count:   2.84803e+06
(500013,474692,499135,401652,498041,474499,0,)
    num_raw_edges:       0 (0,0,0,0,0,0,0,)
    num_waw_edges:       0 (0,0,0,0,0,0,0,)
    num_war_edges:       0 (0,0,0,0,0,0,0,)
    int_num_log:          0 (0,0,0,0,0,0,0,)
    int_debug_get_next:   0 (0,0,0,0,0,0,0,)
    int_debug1:           0 (0,0,0,0,0,0,0,)
    int_debug2:           0 (0,0,0,0,0,0,0,)
    int_debug3:           0 (0,0,0,0,0,0,0,)
    int_debug4:           0 (0,0,0,0,0,0,0,)
    int_debug5:           0 (0,0,0,0,0,0,0,)
    int_debug6:           0 (0,0,0,0,0,0,0,)
    int_debug7:           0 (0,0,0,0,0,0,0,)
    int_debug8:           0 (0,0,0,0,0,0,0,)
    int_debug9:           0 (0,0,0,0,0,0,0,)
    int_debug10:          0 (0,0,0,0,0,0,0,)
    int_psnflush:         0 (0,0,0,0,0,0,0,)
    int_flush_time_interval: 0 (0,0,0,0,0,0,0,)
    int_flush_half_full:  0 (0,0,0,0,0,0,0,)
    int_rec_fail_to_insert: 0 (0,0,0,0,0,0,0,)
    int_num_get_row:      0 (0,0,0,0,0,0,0,)
    int_locktable_volume: 0 (0,0,0,0,0,0,0,)
    int_aux_bytes:        0 (0,0,0,0,0,0,0,)
    int_nonzero:          0 (0,0,0,0,0,0,0,)
    num_log_entries:      0 (0,0,0,0,0,0,0,)

```

```
time_ts_alloc:          0.414599 0.0690998 0.791103% 145.577
(0.0780597,0.0640053,0.0760691,0.0557177,0.0774345,0.0633127,0,)
time_man:               103.657 17.2762 197.79% 36396.8
(17.4497,16.9893,17.4421,17.3147,17.4403,17.021,0,)
time_cleanup:          78.1796 13.0299 149.176% 27451
(13.8762,12.6601,13.8817,11.1327,13.8947,12.7342,0,)
time_txn:              0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_index:            1.28026 0.213376 2.44288% 449.532
(0.206983,0.223236,0.206608,0.202331,0.205339,0.235759,0,)
time_log:              0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_io:               0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_wait_io:          0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_phase1_add_graph: 0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover_txn:      0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_lv_overhead:      0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_state_malloc:     0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_phase1_1:         3.79497 0.632495 7.24125% 1332.52
(0.61578,0.72249,0.548652,0.700306,0.535224,0.672518,0,)
time_phase1_2:         2.59121 0.431869 4.94434% 909.844
(0.396998,0.49119,0.399693,0.414567,0.414911,0.473853,0,)
time_phase2:           5.94461 0.990769 11.343% 2087.31
(0.887513,1.13468,0.92201,0.949424,0.953313,1.09768,0,)
time_phase3:           3.49224 0.58204 6.6636% 1226.22
(0.492516,0.696787,0.497725,0.619835,0.489438,0.695936,0,)
time_phase1_1_raw:     15.5355 2.58925 29.6436% 5454.93
(2.1276,2.55492,2.12508,4.06357,2.10918,2.55515,0,)
time_phase1_2_raw:     0.316563 0.0527605 0.60404% 111.154
(0.0536517,0.0547812,0.0530449,0.0466582,0.0533808,0.0550464,0,)
time_phase2_raw:       0.258381 0.0430635 0.493022% 90.7246
(0.041984,0.0492166,0.0398061,0.0321466,0.041186,0.0540419,0,)
time_phase3_raw:      0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover_full:     0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover1:         0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover2:         0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover3:         0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover4:         0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover5:         0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover6:         0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover7:         0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_recover8:         0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug_get_next:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug0:           0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug1:           0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug2:           0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug3:           0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug4:           0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug5:           2.72727e-08 6.81818e-09 7.80593e-08% 9.57619e-06
(7.27273e-09,0,6.36364e-09,6.36364e-09,7.27273e-09,0,0,)
time_debug6:           0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug7:           0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug8:           0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug9:           0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug10:          0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug11:          19.3175 3.21959 36.8601% 6782.91
```

```

(3.23191,3.2176,3.21231,3.22556,3.1632,3.26696,0,)
time_debug12:          95.3181 15.8863 181.878% 33468.7
(15.8614,15.8935,15.8816,15.8992,15.9306,15.8518,0,)
time_debug13:          0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug14:          0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_debug15:          0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_rec_loop_tryRead: 0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_rec_finding_empty_slot: 0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_silo_validate1:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_silo_validate2:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_silo_validate3:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_silo_validate4:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_silo_validate5:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_silo_validate6:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_silo_validate7:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_silo_validate8:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_silo_validate9:   0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_locktable_get:    16.9353 2.82255 32.3146% 5946.45
(2.30687,2.64078,2.12636,5.14479,2.11503,2.60149,0,)
time_locktable_get_validation: 0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_locktable_release: 0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_get_row_before:    0.847544 0.141257 1.61721% 297.595
(0.123859,0.16625,0.126017,0.141167,0.134412,0.155838,0,)
time_get_row_after:    8.00229 1.33371 15.2693% 2809.82
(1.19266,1.57466,1.35802,0.946161,1.34645,1.58434,0,)
time_log_create:        0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_log_serialLogTxn:  0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_cleanup_1:         0.0261125 0.00435208 0.0498257% 9.16879
(0.00436762,0.00456649,0.00439663,0.0038963,0.00435827,0.00452717,0,)
time_cleanup_2:         0.0449531 0.00749219 0.0857758% 15.7842
(0.00760759,0.00780601,0.00758073,0.00655782,0.00758384,0.00781711,0,)
time_insideSLT1:        0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_insideSLT2:        0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_STLother:          0 -nan -nan% 0 (0,0,0,0,0,0,0,)
time_logging_thread:    0 -nan -nan% 0 (0,0,0,0,0,0,0,)
int_serialLogFail:      0 (0,0,0,0,0,0,0,)
Projected Disk Bandwidth Utilized - avg -nan real -nan
Total time measured 25.6613
Total wall time observed 25.721
Estimated CPU_FREQ is 2.19489
replace 0 tuple
thread 0 use 4967020 tuple
thread 1 use 4647517 tuple
thread 2 use 4815079 tuple
thread 3 use 4167536 tuple
thread 4 use 4410761 tuple
thread 5 use 4267674 tuple
Closing pool file.

```

注意

1. 可通过 STATS_ENABLE 和 STAT_VERBOSE 控制输出信息，前一个控制是否输出，后一个控制统计信息等级（等级为 0 时仅输出部分信息，目前等级为 1）
2. 目前，下面分线程输出时括号里有 7 项，前 6 项表示 6 个线程，最后一项是“log 线程”（?）的信息，不过全是 0
3. 和恢复相关的都没算进去

1. run_time: 单个线程中 在 thread.cpp 里的 run() 函数里的主 while() 循环里执行 run_txn() 及同时进行的一些相关操作的时间的累加和。{ 在 run_txn() 里会再调用提交不同的事务如 payment }

即单个线程的所有事务的执行时间以及其他相关操作时间的累加和

代码位置:

```
124     while (true) {
125         ts_t starttime = get_sys_clock();
126         if (WORKLOAD != TEST) {
127             int trial = 0;
128             if (_abort_buffer_enable) {
129                 m_query = NULL;
130                 while (trial < 2) {
131                     ts_t curr_time = get_sys_clock();
132                     ts_t min_ready_time = UINT64_MAX;
133                     if (_abort_buffer_empty_slots < _abort_buffer_size) {
134                         for (int i = 0; i < _abort_buffer_size; i++) {
135                             if (_abort_buffer[i].query != NULL &&
136                                 curr_time > _abort_buffer[i].ready_time) {
137                                 m_query = _abort_buffer[i].query;
138                                 _abort_buffer[i].query = NULL;
139                                 _abort_buffer_empty_slots++;
140                                 break;
141                             } else if (_abort_buffer_empty_slots == 0 &&
142                                 _abort_buffer[i].ready_time < min_ready_time)
143                                 min_ready_time = _abort_buffer[i].ready_time;
144                         }
145                     }
146                     if (m_query == NULL && _abort_buffer_empty_slots == 0) {
147                         assert(trial == 0);
148                         // M_ASSERT(min_ready_time >= curr_time, "min_ready_time=%ld,
149                         // curr_time=%ld\n", min_ready_time, curr_time);
150                         assert(min_ready_time >= curr_time);
151                         usleep(min_ready_time - curr_time);
152                     } else if (m_query == NULL) {
153                         m_query = query_queue->get_next_query(_thd_id);
154                     #if CC_ALG == WAIT_DIE
155                         m_txn->set_ts(get_next_ts());
156                     #endif
157                     }
158                     if (m_query != NULL)
159                         break;
160                 }
161             } else {
```

```

162         if (rc == RCOK)
163             m_query = query_queue->get_next_query(_thd_id);
164     }
165 }
166 INC_STATS(_thd_id, time_query, get_sys_clock() - starttime);
167 m_txn->set_start_time(get_sys_clock());
168 m_txn->abort_cnt = 0;
169 // #if CC_ALG == VLL
170 //     _wl->get_txn_man(m_txn, this);
171 // #endif
172 m_txn->set_txn_id(get_thd_id() + thd_txn_id * g_thread_cnt); // qfs
173 thd_txn_id++;
174
175 if ((CC_ALG == HSTORE && !HSTORE_LOCAL_TS) || CC_ALG == MVCC ||
176     CC_ALG == HEKATON || CC_ALG == TIMESTAMP)
177     m_txn->set_ts(get_next_ts());
178
179 rc = RCOK;
180 #if CC_ALG == HSTORE
181     rc = part_lock_man.lock(m_txn, m_query->part_to_access, m_query->part_num)
182 #elif CC_ALG == VLL
183     vll_man.vllMainLoop(m_txn, m_query);
184 #elif CC_ALG == MVCC || CC_ALG == HEKATON
185     glob_manager->add_ts(get_thd_id(), m_txn->get_ts());
186 #elif CC_ALG == OCC
187     // In the original OCC paper, start_ts only reads the current ts without
188     // advancing it. But we advance the global ts here to simplify the
189     // implementation. However, the final results should be the same.
190     m_txn->start_ts = get_next_ts();
191 #endif
192 if (rc == RCOK) {
193     #if CC_ALG != VLL
194         rc = m_txn->run_txn(m_query);
195     #endif
196     #if CC_ALG == HSTORE
197         part_lock_man.unlock(m_txn, m_query->part_to_access, m_query->part_num);
198     #endif
199 }
200 if (rc == Abort) {
201     // cout << m_txn->get_txn_id() << " Aborted" << endl;
202     uint64_t penalty = 0;
203     if (ABORT_PENALTY != 0) {
204         double r;
205         // drand48_r(&buffer, &r);
206         r = erand48(buffer);
207         penalty = r * ABORT_PENALTY;
208     }
209     if (!_abort_buffer_enable)
210         usleep(penalty / 1000);
211     else {
212         assert(_abort_buffer_empty_slots > 0);
213         for (int i = 0; i < _abort_buffer_size; i++) {
214             if (_abort_buffer[i].query == NULL) {
215                 _abort_buffer[i].query = m_query;
216                 _abort_buffer[i].ready_time = get_sys_clock() + penalty;
217                 _abort_buffer_empty_slots--;

```

```

218         break;
219     }
220 }
221 }
222 }
223
224     ts_t endtime = get_sys_clock();
225     uint64_t timespan = endtime - starttime;
226     INC_FLOAT_STATS_V0(run_time, timespan);

```

2. total_run_time: 6 个线程 run_time 的和
3. num_commits: 单个线程提交的总事务数
4. total_num_commits: 6 个线程提交的总事务数

以下公式里忽略掉单位相关的乘数

1. Throughput: 吞吐量

公式: $\text{total_num_commits} / (\text{total_run_time} / 6)$

2. MaxThr: 不知道是啥, 跟吞吐量差不多

公式: $\text{total_num_commits} / \text{max_run_time}$

3. PerThdThr: 线程平均吞吐量?

公式: 6 个线程的 $(\text{num_commits} / \text{run_time})$ 的和

4. run_time: 最前面的是 total_run_time, 后面的是 6 个线程各自的 run_time。以下 5, 6, 7, 8 格式与这里相同。

5. latency: start 是在 thread.cpp 的主循环里 run_txn() 函数前的 set_start_time(), end 是在 cleanup() 里, latency=end-start。latency 比 run_time 大的原因: run_time 在输出前会除以 CPU_FREQ(=2.2), 而 latency 不会; 实际上二者确实是大约相差两倍。

6. num_commits: 见前

7. num_aborts: 线程 abort 事务数

8. num_latency_count: 计算 latency 的次数。单个线程中, latency 每累加、改变一次, 该变量加 1

9. time_ts_alloc: 线程里获得全局时间戳耗费的时间的累加 (比如包括互斥锁的开销)。

第一项是 6 个线程的 total 值, 第二项是去掉 0 值的 6 个线程平均值 nonzero_avg; 第三项是一个百分数: $\text{nonzero_avg} / (\text{total_run_time} / 6)$; 第四项是 $\text{total} / \text{total_num_commits}$; 最后面是 6 个线程的各自的值。下面的各个变量格式也是这样。

10. time_man: 线程里 get_row() 的时间的累加和 finish() 的时间的累加的和。 (目前 finish() 里基本上只有 cleanup())

11. time_cleanup: 线程里 finish() 的时间的累加。 (目前 finish() 里基本上只有 cleanup())

12. time_index: 线程里 查索引(index_read()函数)的时间的累加。

13. time_phase1_1: new_order 事务里第 1-1 部分 SQL 语句耗费的时间 加上 cleanup()里见下图的部分感觉主要是 return_row()函数耗费的时间 的和

```
system > C++ txn.cpp
400 uint64_t starttime = get_sys_clock();
401 // start to release the locks
402 #if CC_ALG != SILO // updating the data is already handled in silo_validate
403     for (int rid = row_cnt - 1; rid >= 0; rid--) {
404         row_t *orig_r = accesses[rid]->orig_row;
405         access_t type = accesses[rid]->type;
406         if (type == WR && rc == Abort)
407             type = XP; // means we need to roll back the data value
408
409         #if (CC_ALG == NO_WAIT || CC_ALG == DL_DETECT) &&
410             ISOLATION_LEVEL == REPEATABLE_READ
411             if (type == RD) {
412                 accesses[rid]->data = NULL;
413                 continue;
414             }
415         #endif
416
417         char *newdata;
418         if (ROLL_BACK && type == XP &&
419             (CC_ALG == DL_DETECT || CC_ALG == NO_WAIT || CC_ALG == WAIT_DIE)) {
420             #if NEW_ROW
421                 newdata = ((row_t *)accesses[rid]->orig_data)->tupple_data->data;
422             #else
423                 newdata = ((row_t *)accesses[rid]->orig_data)
424                     ->data; // fixed a bug left from the original code base
425             #endif
426
427             } else {
428                 newdata = accesses[rid]->data;
429             }
430         #if USE_LOCKTABLE
431             LockTable &lt = LockTable::getInstance();
432             uint64_t current_time = get_sys_clock();
433             // assert((uint64_t)newdata != 0);
434             #if LOG_ALGORITHM == LOG_TAURUS
435                 lt.release_lock(orig_r, type, this, newdata, lsn_vector, NULL, rc);
436             #elif LOG_ALGORITHM == LOG_SERIAL
437                 lt.release_lock(orig_r, type, this, newdata, NULL, &_max_lsn, rc);
438             #else
439                 lt.release_lock(orig_r, type, this, newdata, NULL, NULL, rc);
440             #endif
441             INC_INT_STATS(time_locktable_release, get_sys_clock() - current_time);
442         #else
443             orig_r->return_row(type, this, newdata, accesses[rid]->row_data, rc);
444         #endif
445         #if CC_ALG != TICTOC && CC_ALG != SILO
446             accesses[rid]->data = NULL; // will not need this any more
447         #endif
448     }
449     #if VERBOSE_LEVEL > 0
450         stringstream ssk;
451         ssk << GET_THD_ID << " finishes" << endl << endl;
452         cout << ssk.str();
453     #endif
454
455 #endif
456 uint64_t cleanup_1_begin = get_sys_clock();
457 INC_INT_STATS(time_phase1_1, cleanup_1_begin - starttime);
458 if (rc == Abort) { // remove inserted rows
```

为啥是这两部分，没啥规律啊

14. time_phase1_2: new_order 事务里第 1-2 部分 SQL 语句耗费的时间 加上 cleanup()里一部分和 log 相关的代码（我感觉目前 log_no 时实际上这一部分什么也没执行）耗费的时间的和。

15. time_phase2: new_order 事务里第 2 部分 SQL 语句耗费的时间。

- 16. time_phase3: new_order 事务里第 3 部分 SQL 语句耗费的时间。
- 17. time_phase1_1_raw: new_order 事务里第 3 部分 SQL 语句后面一部分语句耗费的时间。
- 18. time_phase1_2_raw: new_order 事务里再往后,,,,,
- 19. time_phase2_raw: new_order 事务里再往后,,,,,
- 20. time_debug5: 在 log.cpp 里几个函数里耗费的时间(不清楚这几个函数的用途和是否调用) 加上 try_commit_txn()里耗费的时间
- 21. time_debug11:在 tpcc_txn.cpp 的 run_txn () 里执行 run_payment()函数耗费的时间(见下图) 加上 log.cpp 里几个函数里耗费的时间(不清楚这几个函数的用途和是否调用)

```
26  RC tpcc_txn_man::run_txn(base_query *query, bool rec) {
27      _query = (tpcc_query *)query;
28      RC rc = RCOK;
29      // cout << _query->type << endl;
30      uint64_t starttime = get_sys_clock();
31      switch (_query->type) {
32      case TPCC_PAYMENT:
33          rc = run_payment(_query);
34          COMPILER_BARRIER
35          INC_INT_STATS(time_debug11, get_sys_clock() - starttime);
36          break;
```

- 22. time_debug12:在 tpcc_txn.cpp 的 run_txn () 里执行 run_new_order()函数耗费的时间(见下图) 加上 log.cpp 里几个函数里耗费的时间(不清楚这几个函数的用途和是否调用)

```
26  RC tpcc_txn_man::run_txn(base_query *query, bool rec) {
27      _query = (tpcc_query *)query;
28      RC rc = RCOK;
29      // cout << _query->type << endl;
30      uint64_t starttime = get_sys_clock();
31      switch (_query->type) {
32      case TPCC_PAYMENT:
33          rc = run_payment(_query);
34          COMPILER_BARRIER
35          INC_INT_STATS(time_debug11, get_sys_clock() - starttime);
36          break;
37      case TPCC_NEW_ORDER:
38          rc = run_new_order(_query);
39          COMPILER_BARRIER
40          INC_INT_STATS(time_debug12, get_sys_clock() - starttime);
41          break;
```

- 23. time_locktable_get: 见下图, 基本上是只有 810 行 get_row()里再调用另一个 get_row()函数所耗费的时间


```

790     uint64_t right_before_get = get_sys_clock();
791     INC_INT_STATS_V0(time_get_row_before, right_before_get - starttime);
792     // if(row_cnt == 19)
793     // printf("Encounter 19, before get_row as %" PRIu64 " , type %d\n",
794     // (uint64_t)row, type);
795     #if USE_LOCKTABLE && CC_ALG != SILO
796         LockTable &ltlt = LockTable::getInstance();
797     #endif
798     #if LOG_ALGORITHM == LOG_TAURUS
799         rc = lt.get_row(row, type, this, accesses[row_cnt]->data, lsn_vector, NULL);
800     #elif LOG_ALGORITHM == LOG_SERIAL
801         rc = lt.get_row(row, type, this, accesses[row_cnt]->data, NULL,
802         | | | | | | | | &_max_lsn); //, true, 0, true);
803     #else
804         rc = lt.get_row(row, type, this, accesses[row_cnt]->data, NULL, NULL);
805     #endif
806
807     #else
808         // if (accesses[row_cnt]->row_data == NULL)
809         // cout << "qfs::" << 1 << " " << type << "!!\n";
810         rc = row->get_row(type, this, accesses[row_cnt]->data,
811         | | | | | | | | accesses[row_cnt]->row_data);
812     #endif
813
814     #if LOG_ALGORITHM == LOG_TAURUS
815         partition_accesses_cnt[logPartition((uint64_t)row)]++;
816     #endif
817
818     uint64_t starttime2 = get_sys_clock();
819     INC_INT_STATS_V0(time_locktable_get, starttime2 - right_before_get);

```

24. time_get_row_before: 见下图，get_row()里这一部分所耗费的时间。

```

769     uint64_t starttime = get_sys_clock();
770     RC rc = RCOK;
771     if (row_cnt == num_accesses_alloc) {
772         // accesses[row_cnt] == NULL // bad
773         Access *access = (Access *)MALLOC(sizeof(Access), GET_THD_ID);
774         accesses[row_cnt] = access;
775     }
776     #if (CC_ALG == SILO || CC_ALG == TICTOC)
777         access->data = new char[MAX_TUPLE_SIZE];
778         // #if LOG_ALGORITHM == LOG_TAURUS
779         // access->orig_data = (char*) row;
780         // #else
781         access->orig_data = NULL;
782     // #endif
783     #elif (CC_ALG == DL_DETECT || CC_ALG == NO_WAIT || CC_ALG == WAIT_DIE)
784         accesses[row_cnt]->orig_data = (char *)MALLOC(sizeof(row_t), GET_THD_ID);
785         ((row_t *) (accesses[row_cnt]->orig_data))->init(MAX_TUPLE_SIZE);
786     #endif
787
788     num_accesses_alloc++;
789 }
790     uint64_t right_before_get = get_sys_clock();
791     INC_INT_STATS_V0(time_get_row_after, right_before_get - starttime);

```

25. time_get_row_after: get_row()里从 time_locktable_get 结束的时间点到函数最后所耗费的时间。

26. time_cleanup_1: cleanup()里这一部分代码耗费的时间。

```

456     uint64_t cleanup_1_begin = get_sys_clock();
457     INC_INT_STATS(time_phase1_1, cleanup_1_begin - starttime);
458     if (rc == Abort) { // remove inserted rows.
459         for (UInt32 i = 0; i < insert_cnt; i++) {
460             row_t *row = insert_rows[i];
461             assert(g_part_alloc == false);
462             #if CC_ALG != HSTORE && CC_ALG != OCC
463                 mem_allocator.free(row->manager, 0);
464             #endif
465             row->free_row();
466             mem_allocator.free(row, sizeof(row));
467         }
468     }
469     uint64_t cleanup_1_end = get_sys_clock();
470     INC_INT_STATS(time_cleanup_1, cleanup_1_end - cleanup_1_begin);

```

27. time_cleanup_2: cleanup()里这一部分代码耗费的时间。

```

614     uint64_t cleanup2_begin = get_sys_clock();
615     INC_INT_STATS(time_phase1_2, cleanup2_begin - cleanup_1_end);
616     try_commit_txn(); // no need to try_commit_txn if abort
617     #else // LOG_ALGORITHM == LOG_NO
618         uint64_t cleanup2_begin = get_sys_clock();
619         INC_INT_STATS_V0(num_latency_count, 1);
620         INC_FLOAT_STATS_V0(latency, get_sys_clock() - _txn_start_time);
621     #endif
622
623     #if LOG_ALGORITHM == LOG_PLOVER
624         memset(_log_entry_sizes, 0, sizeof(uint32_t) * g_num_logger);
625     #else
626         _log_entry_size = 0;
627     #endif
628     row_cnt = 0;
629     wr_cnt = 0;
630     insert_cnt = 0;
631     #if LOG_ALGORITHM == LOG_PARALLEL
632         _num_raw_preds = 0;
633         _num_waw_preds = 0;
634     #elif LOG_ALGORITHM == LOG_SERIAL
635         _max_lsn = 0;
636     #elif LOG_ALGORITHM == LOG_TAURUS
637         _max_lsn = 0;
638         memset(partition_accesses_cnt, 0, sizeof(uint64_t) * g_num_logger);
639         memset(lsn_vector, 0, sizeof(lsnType) * g_num_logger);
640     #endif
641     #if CC_ALG == DL_DETECT
642         dl_detector.clear_dep(get_txn_id());
643     #endif
644     INC_INT_STATS(time_cleanup_2, get_sys_clock() - cleanup2_begin);

```